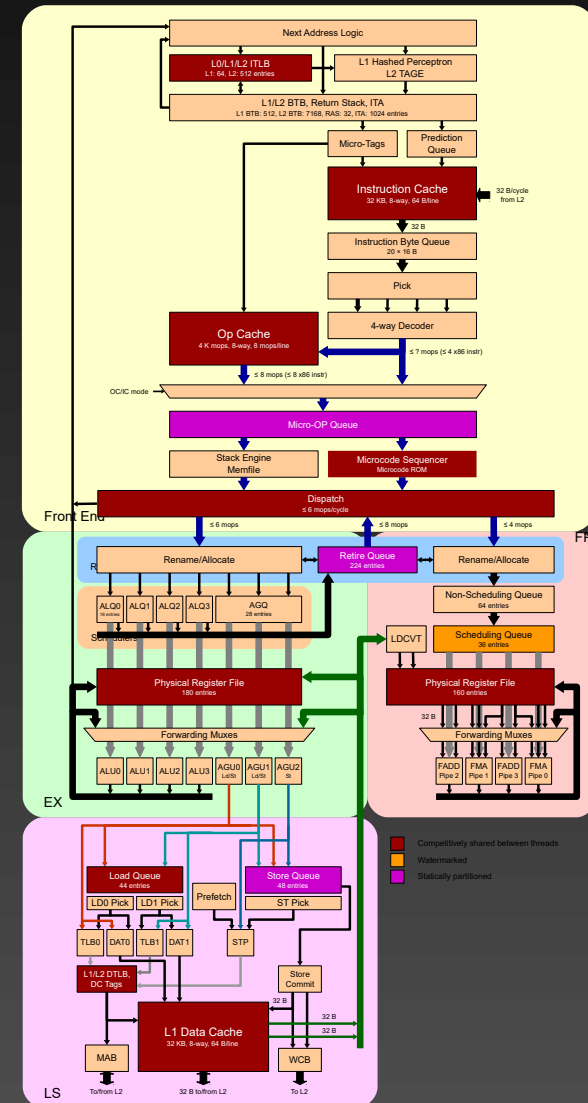



Concurrent vs. Parallel vs. Distributed

- Concurrent
 - different parts of the program running in single system at the same time, can communicate
- Parallel
 - single computation divided to smaller and same parts, running concurrently
- Distributed
 - concurrent execution on different computers

HW Parallelism

- Registers
 - SSE = 128b, AVX(2) = 256b, AVX512, ...
 - parts of 2^n bytes
- Instructions
 - more ALUs
 - Out Of Order (OOO) execution
- SMT = Simultaneous Multi-Threading
 - more threads (2-16) in single core
 - better ALU utilization
- SMP – Symmetric Multi-Processing
- Cluster – computers + fabric (FAST interconnect)
- Grid, Cloud – servers + common network



Shared Data

- Possible **Race Condition**
 - Conflict over a resource without coordination
 - Bad things happen as a result – Undefined Behaviour
- **Data Race Definition**
 - 1) Two or more threads access the same memory
 - 2) At least one access is a write
 - 3) The threads do not synchronize with each other
- Synchronisation (locking) necessary
- Locks
 - mutex, critical section, synchronized methods, semaphor, ...

Implicit vs. Explicit Parallelism

- **Implicit**
 - Automatic parallelisation by compiler
 - both instruction level and pieces of source code (usually for loops)
 - Precompiled libraries (OpenCV)
 - Implicitly parallel programming languages
 - LabView, Matlab (1:N)
 - No effort to splitting, comm, sync
 - Smaller control over runtime, smaller efficiency, overhead is hidden

Compiler Based Implicit Parallelism

- Modern compilers perform auto-vectorisation
 - Compiler Explorer <https://godbolt.org/>

The screenshot displays the Compiler Explorer interface. On the left, the C++ source code for a function `testFunction` is shown. The code calculates the sum of an array of integers. The middle pane shows the assembly output for `x86-64 gcc 13.2` with optimization level `-O2`. The assembly includes instructions for testing, comparing, and adding, with labels `.L3` and `.L4`. The right pane shows the assembly output for `x86-64 gcc 13.2` with optimization level `-O3`. This assembly includes SIMD instructions like `paddq`, `movdqa`, and `psrldq`, indicating auto-vectorization. A tooltip is visible over the SIMD instructions, explaining that they perform a SIMD add of packed integers. The bottom status bar shows the compilation time and line count for both versions.

```
1 int testFunction(int* input, int length) {
2     int sum = 0;
3     for (int i = 0; i < length; ++i) {
4         sum += input[i];
5     }
6     return sum;
7 }
8
```

x86-64 gcc 13.2 (Editor #1) -O2

```
1 testFunction(int*, int):
2     test    esi, esi
3     jle     .L4
4     movsx   rsi, esi
5     xor     eax, eax
6     lea     rdx, [rdi+rsi*4]
7 .L3:
8     add     eax, DWORD PTR [rdi]
9     add     rdi, 4
10    cmp     rdi, rdx
11    jne     .L3
12    ret
13 .L4:
14    xor     eax, eax
15    ret
```

x86-64 gcc 13.2 (Editor #1) -O3

```
14 .L4:
15     movdqa  xmm2, xmmword ptr [rax]
16     pshufb  xmm2, xmm2, 0x55
17     movdqa  xmm1, xmm2
18     movdqa  xmm0, xmm1
19     psrldq  xmm1, 4
20     paddq   xmm0, xmm1
21     movd    eax, xmm0
22     test    sil, 3
23     je      .L11
24     paddq   xmm0, xmm1
25     movdqa  xmm1, xmm0
26     psrldq  xmm1, 4
27     paddq   xmm0, xmm1
28     movd    eax, xmm0
29     test    sil, 3
30     je      .L11
31 .L3:
32     movsx   rdi, edx
33     lea     r8, [0+rdi*4]
```

Performs a SIMD add of the packed integers from the source operand (second operand) and the destination operand (first operand), and stores the packed integer results in the destination operand. See Figure 9-4 in the Intel® 64 and IA-32 Architectures Software Developer's Manual, Volume 1, for an illustration of a SIMD operation. Overflow is handled with wraparound, as described in the following paragraphs. More information available in the context menu.

Output (0/0) x86-64 gcc 13.2 - 612ms (6859B) ~427 lines filtered

Compiler License

Configuration: Release

Platform: Active(x64)

Configuration Manager...

- Configuration Properties
- General
- Advanced
- Debugging
- VC++ Directories
- C/C++
- General
- Optimization
- Preprocessor
- Code Generation
- Language
- Precompiled Headers
- Output Files
- Browse Information
- External Includes
- Advanced
- All Options
- Command Line
- Linker
- Manifest Tool
- XML Document Generator
- Browse Information
- Build Events
- Custom Build Step
- Code Analysis

Enable String Pooling	
Enable Minimal Rebuild	No (/Gm-)
Enable C++ Exceptions	Yes (/EHsc)
Smaller Type Check	No
Basic Runtime Checks	Default
Runtime Library	Multi-threaded DLL (/MD)
Struct Member Alignment	Default
Security Check	Enable Security Check (/GS)
Control Flow Guard	
Enable Function-Level Linking	Yes (/Gy)
Enable Parallel Code Generation	
Enable Enhanced Instruction Set	Not Set
Floating Point Model	CPU extension requirements ARMv8.0-A (ARM64) (/arch:armv8.0)
Enable Floating Point Exceptions	CPU extension requirements ARMv8.1-A (ARM64) (/arch:armv8.1)
Create Hotpatchable Image	CPU extension requirements ARMv8.2-A (ARM64) (/arch:armv8.2)
Spectre Mitigation	CPU extension requirements ARMv8.3-A (ARM64) (/arch:armv8.3)
Enable Intel JCC Erratum Mitigation	CPU extension requirements ARMv8.4-A (ARM64) (/arch:armv8.4)
Enable EH Continuation Metadata	CPU extension requirements ARMv8.5-A (ARM64) (/arch:armv8.5)
Enable Signed Returns	CPU extension requirements ARMv8.6-A (ARM64) (/arch:armv8.6)
	CPU extension requirements ARMv8.7-A (ARM64) (/arch:armv8.7)
	CPU extension requirements ARMv8.8-A (ARM64) (/arch:armv8.8)
	No Enhanced Instructions (/arch:IA32)
	Not Set

Enable Enhanced Instruction Set
Enable use of instructions found on processors that support enhanced instruction sets. If no option is specified, the compiler will use instructions found on processors that support SSE2. Use of enhanced instructions can be disabled with /arch:IA32. (/arch:SSE, /arch:SSE2, /arch:AVX, /arch:AVX2, /arch:A...

C++ Parallel Algorithms

- Let the compiler do all the hard work for you (min. C++ 17)
 - Add **execution policy** as first argument to **algorithm** call
- Execution policies
 - **std::execution::seq**
 - Run sequentially, no parallelism
 - **std::execution::par**, **std::execution::par_unseq**
 - Request to compiler to run in parallel
 - Promise by user that code is safe to run in parallel; no data races
 - **std::execution::unseq**, **std::execution::par_unseq**
 - Request to compiler to vectorize
 - Promise by user that code is safe to vectorize; no data races or locks

C++ Parallel Algorithms Examples

```
std::sort(items.begin(), items.end(), compare_by_price);
```

→

```
std::sort(std::execution::par, items.begin (), items.end(), compare_by_price);
```

```
std::fill(v.begin(), v.end(), 1);
```

→

```
std::fill(std::execution::par_unseq, v.begin(), v.end (), 1);
```


Some of C++ Parallel Algorithms

`for_each, for_each_n, transform`

`find, find_if, find_end, search, count, any_of, adjacent_find`

`copy, copy_if, move, fill, replace, generate, rotate`

`sort, stable_sort, partial_sort, nth_element, is_sorted`


`reduce, transform_reduce, inclusive_scan, exclusive_scan`

Implicit vs. Explicit Parallelism


- **Explicit**
 - Precise control of concurrency, sync, comm using compiler directives, function calls etc.
 - overhead is visible and controllable
 - Directives
 - code splitting
 - synchronisation
 - communication
 - Full user control, higher efficiency possible
 - thread API, OpenMP, MPI

Data vs Task Parallelism

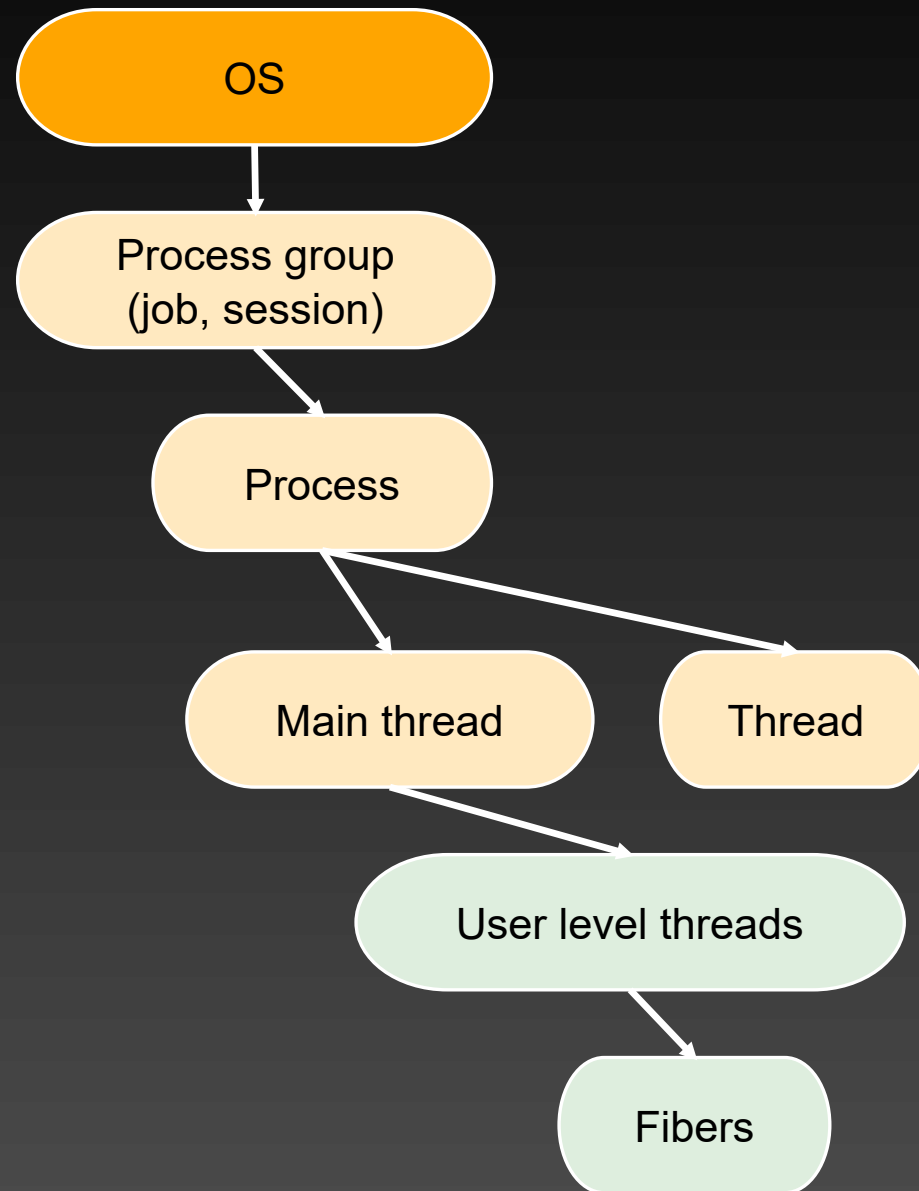
- Data-parallel
 - data are distributed
 - thread code (nearly) same



- Task-parallel
 - code is distributed

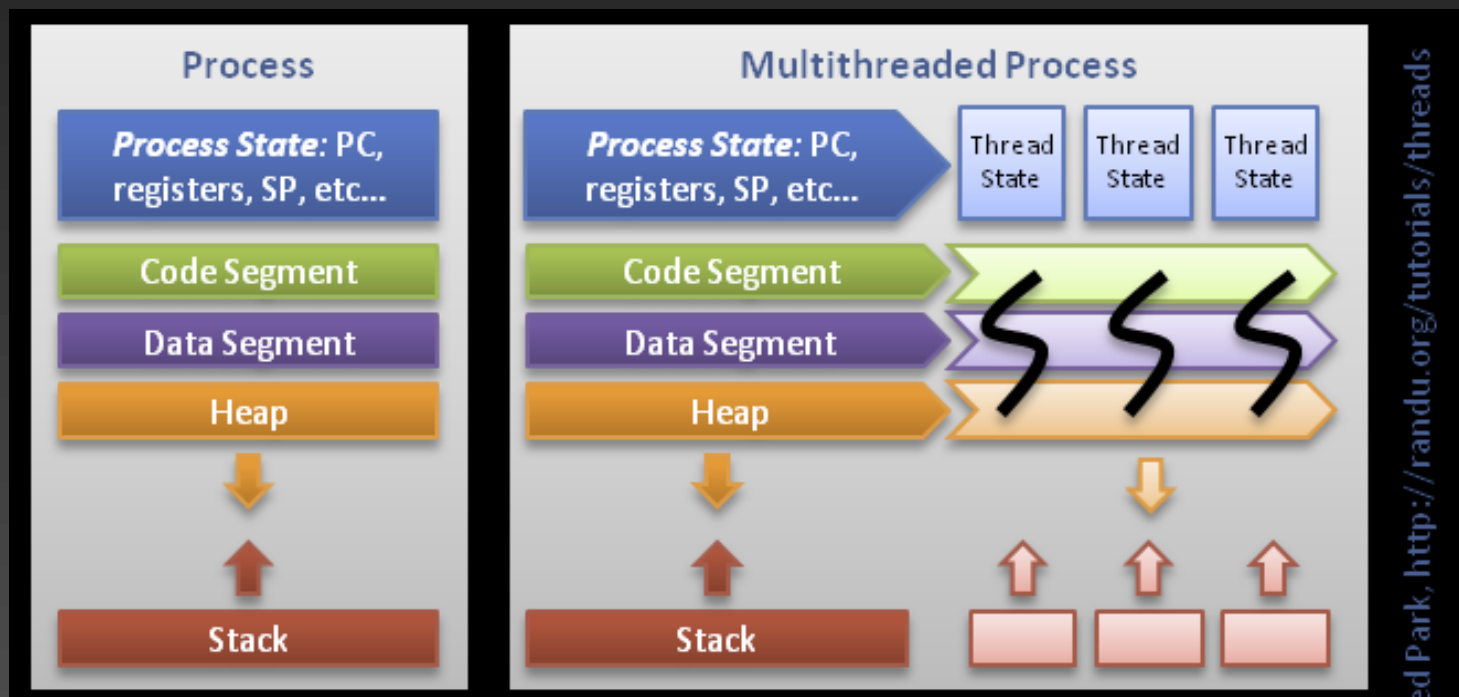


Task Hierarchy in OS



Process vs. Thread

- Process
 - Virtual memory, privileges, code, PID, priority
 - at least one thread
- Thread
 - memory is shared for all threads
 - thread local: only stack, registers, thread ID



Splitting APP to Threads

- Usually data and task parallelism
- Design patterns
 - Master/Slave + thread pool
 - master thread scatter and gather data + control others, does NOT compute itself
 - workers (slaves) threads usually created in advance
 - Equal threads
 - master also works
 - Pipeline
 - task parallel, each threads does different task
 - problems if one stage is slower


Types of Parallelism

- Fine grained
 - frequent comm and sync
 - small data blocks after shot execution
 - very latency sensitive
- Coarse grained
 - occasional communication
 - sync necessary, but not latency sensitive
- Embarrassingly parallel
 - completely independent tasks, zero comm
 - e.g. repeated run with different cmd line args

Synchronisation

- MUTEX = MUTual EXclusion
 - lock used for serialisation of thread access to resources
- Critical section
 - code between mutex locking and unlocking
 - guaranteed to be executed by only single thread at once
 - serial time
 - must be as small as possible
- Atomic operations
 - simple operations, guaranteed by hardware or library to be correct without explicit mutex (e.g. ++)

Parallel vs. Serial Time



Two independent parts **A** **B**

Original process



Make **B** 5x faster



Make **A** 2x faster



Safety

- Dangerous operation considering parallel execution
 - uncontrolled access to globals (variables and heap)
 - saving function state to global variables
 - global resource (de)allocations (files, sockets, graphics, ...)
 - indirect access to data using pointers and references
 - visible side effects (modifications of volatile variables)
- Safe strategy
 - use only local variables (stack) and `thread_local` variables
 - code depends only on function arguments, value passing
 - all functions and subfunctions are re-entrant

Synchronisation primitives I

- Mutex – lock (locked vs. unlocked)
 - `std::mutex`
`void lock(); bool try_lock(); void unlock();`
- Barrier (C++ 20 `std::barrier`, `std::latch`)
 - position in code, where execution of a thread is paused until all threads will arrive
- Join (fork-join)
 - gather results and exit status from all threads – join will terminate thread

Synchronisation primitives II

- Conditional variable – `std::conditional_variable`
 - call `wait()` for variable → thread put to sleep
 - HW watching for write into variable → wake up
 - have to check for wake-up reason

```
lock( mutex_x );  
while ( not wake_me ) { sleep(cond_var, mutex_x); }  
unlock( mutex_x );
```

- Semaphore – `std::counting_semaphore` (C++ 20)
 - special case: binary = mutex
 - maintains internal counter, – set to N (resources), thread enters → `acquire()` = N--
thread exits → `release()` = N++
on N=0 → block and wait

Threads

- Each program has one main thread, created by OS
 - Other threads created explicitly
 - Each thread can start more threads
- Thread is immediately ready to run
 - It can be started by OS scheduler before parent thread returns from create function
 - All data necessary for thread must be prepared BEFORE calling create

Thread properties

- Detached threads

 - `void detach();`

 - Can not be joined with master by `join()`
 - Run on background, saves app resources
 - Master thread **can not** obtain thread exit code – no way to detect error

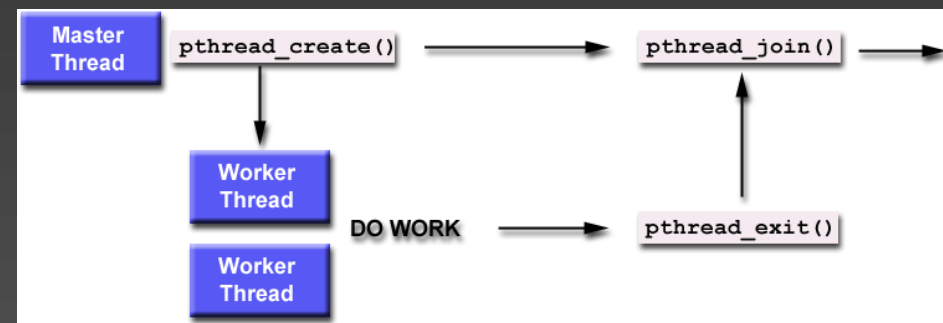
 - programmer must create some other way to send data (status, result) from detached thread to main thread

- `std::jthread` (C++ 20)

 - automatically rejoins on destruction
 - destructor called automatically at the end of the scope
 - no need to call `join()` manually

Terminating thread

- Thread can be terminated
 - by running to the end of code
 - by ending parent thread execution by `exit()`, `abort()`, `return`, `std::terminate()` ...
 - by ending master thread other than `return` (`kill`, `exit`, `abort`...)
- Thread after termination
 - process resources (`fd`, `IPC`, `mutex`, ...) created (opened) in thread are **NOT** closed (deallocated) – global resources
 - heap data referenced only from thread must be released before `exit` – otherwise memory leak (system will release all resources on process termination, not thread termination)
- Join
 - Thread structures (status code) are cleared after thread join
 - Blocking call - wait for thread to finish
 - Join necessary if we want to know exit status code




Producer – Consumer

- Thread P produces data, C consumes
- Possible solution
 - shared data storage + mutex
- Better
 - common queue
 - counting semaphore, P increases, C decreases
 - at zero C can be put to sleep
 - problem: more C or P – removing and inserting is not atomic, PxC resource overwriting
- Best
 - conditional variables
 - any amount of C a P, single storage, no busy wait

Data Structures


ConcurrentStack



Consumers and producers
compete for the top of the stack

Data Structures

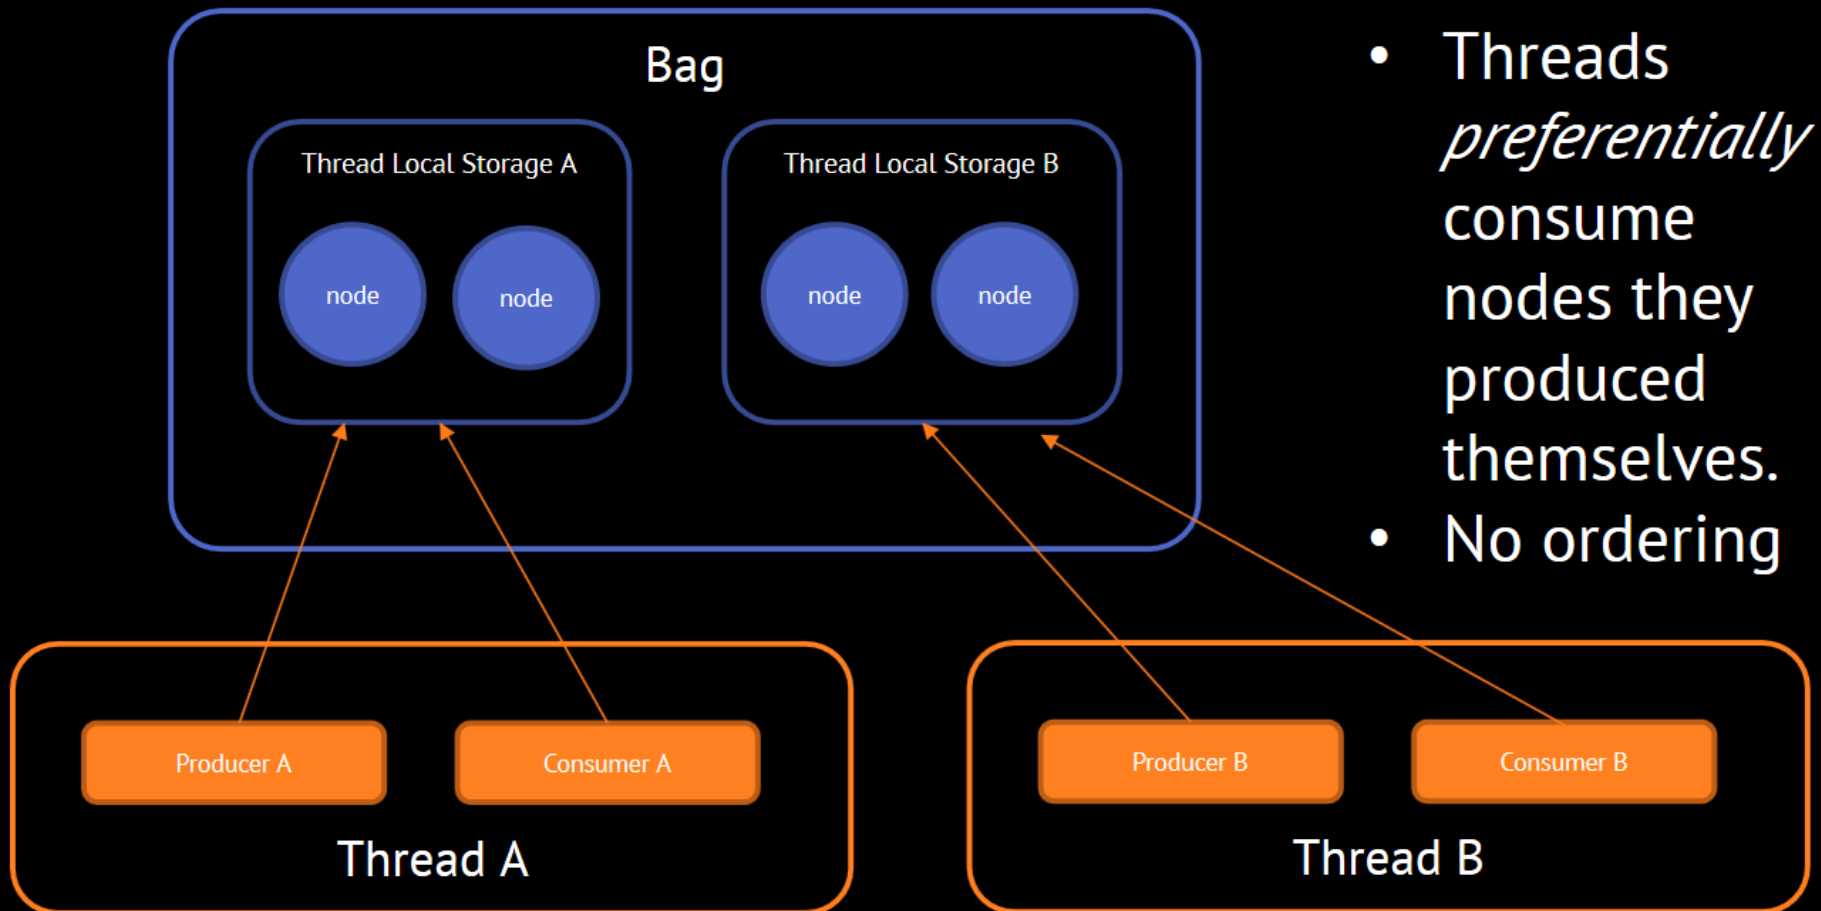
Concurrent Queue



Consumers compete for the head.
Producers compete for the tail.
Both compete for the same node if the queue is empty.

Data Structures

ConcurrentBag



HW Threads in C++

- header <thread>
- How many HW threads are supported by my CPU? (including SMT)

```
#include <iostream>
#include <thread>
#include <chrono>

int main(int argc, char** argv) {
    const unsigned int hw = std::thread::hardware_concurrency();

    std::cout << "Wait a moment...\n";           // End-Of-Line without flush (faster)

    std::this_thread::sleep_for(std::chrono::milliseconds(25));
    std::this_thread::sleep_for(std::chrono::seconds(3));

    std::cout << "Got HW threads: " << hw << std::endl; // End-Of-Line with implicit flush (safer, slower)
    return EXIT_SUCCESS;
}
```

Synchronization by Fork & join

- create single thread vs. multiple threads, pass a parameter

```
#include <iostream>
#include <thread>

void thread_code(void) {
    std::cout << "Hello world from thread: ID=" <<
        std::this_thread::get_id() << "\n";
}

int main(int argc, char** argv) {
    std::thread my_thread(thread_code); // FORK
    // now thread_code() is running in separate thread...

    my_thread.join();                // JOIN

    return EXIT_SUCCESS;
}
```

```
#include <iostream>
#include <vector>
#include <thread>

static const int num_threads = 10;

void thread_code(const int i) {
    std::cout << "Hello world from thread [" << i <<
        "]" : ID=" << std::this_thread::get_id() <<
        "\n";
}

int main(int argc, char** argv) {
    std::vector<std::thread> threads;

    threads.resize(num_threads);
    for (int i = 0; i < num_threads; ++i) {
        threads[i] = std::thread(thread_code, i);
        // additional parameteres passed to thread function
    }

    for (int i = 0; i < 10; ++i) {
        threads[i].join();
    }

    return EXIT_SUCCESS;
}
```

Thread Identification

- Identify main thread & execute code accordingly

```
#include <iostream>
#include <thread>

std::thread::id main_thread_id = std::this_thread::get_id();

void am_i_main(void)
{
    if (main_thread_id == std::this_thread::get_id())
        std::cout << "This is the main thread.\n";
    else
        std::cout << "This is not the main thread.\n";
}

void thread_code(void) {
    std::cout << "Hello world from thread: " << std::this_thread::get_id() << ". ";
    am_i_main();
}

int main(int argc, char** argv) {
    std::thread my_thread(thread_code);

    am_i_main();

    my_thread.join();

    return EXIT_SUCCESS;
}
```

```
This is the main thread.
Hello world from thread: 3832. This is not the main thread.
```

std::mutex DO NOT USE

- hardware synchronised inter-thread communication
 - in critical section, more complex operations are possible

```
#include <iostream>
#include <vector>
#include <chrono>
#include <thread>
#include <mutex> // std::mutex, std::scoped_lock

static std::mutex my_mutex;

static const int num_threads = 10;

void thread_code(const int tid, int& result) {
    std::this_thread::sleep_for(std::chrono::seconds(1));
    std::thread::id this_id = std::this_thread::get_id();

    // try to move lock BELOW printing
    my_mutex.lock();
    std::cout << "I am " << tid <<
        " with id " << this_id << std::endl;

    result += 1;
    my_mutex.unlock();
}

int main(int argc, char** argv) {
    std::vector<std::thread> threads;
    int result = 0;

    threads.resize(num_threads);
    for (int i = 0; i < num_threads; ++i) {
        threads[i] = std::thread(thread_code, i,
std::ref(result));
    }

    for (int i = 0; i < num_threads; ++i) {
        threads[i].join();
    }

    std::cout << "Result: " << result << std::endl;
    return EXIT_SUCCESS;
}
```

some lock_guard variant USE

```
#include <iostream>
#include <vector>
#include <chrono>
#include <thread>
#include <mutex> // std::mutex, std::scoped_lock

static std::mutex my_mutex;

static const int num_threads = 10;

void thread_code(const int tid, int& result) {
    std::this_thread::sleep_for(std::chrono::seconds(1));
    std::thread::id this_id = std::this_thread::get_id();
    std::scoped_lock lock(my_mutex);

    std::cout << "I am " << tid <<
        " with id " << this_id << std::endl;

    result += 1;
}

int main(int argc, char** argv) {
    std::vector<std::thread> threads;
    int result = 0;

    threads.resize(num_threads);
    for (int i = 0; i < num_threads; ++i) {
        threads[i] = std::thread(thread_code, i,
std::ref(result));
    }

    for (int i = 0; i < num_threads; ++i) {
        threads[i].join();
    }

    std::cout << "Result: " << result << std::endl;

    return EXIT_SUCCESS;
}
```

Lock Guards

- RAII wrapper around mutexes
 - Constructor calls lock()
 - Destructor calls unlock()
- Guarantee that the mutex is always released

`std::scoped_lock` (C++ 17)

Lock Guards

`std::scoped_lock` (C++ 17)

- Constructor takes one or more mutexes
 - Calls `lock()` on each of the mutexes
- Destructor calls `unlock()` on each of the mutexes
- Not copyable or movable
- No member functions or other operations

```
std::scoped_lock my_lock(mutex_a , mutex_b);
```

- `std::scoped_lock` is useful for **avoiding deadlock**
- If given multiple mutexes, always locks them in the same order

Lock Guards

`std::unique_lock`

- Owns a mutex
- Destructor calls `unlock()` on the mutex if the mutex is locked
- Has `lock()`, `unlock()`, and several other member functions
- Movable, but not copyable
- Useful when you need more control over when the mutex is locked

Atomic

- hardware synchronised inter-thread communication
 - use for smallest possible amounts of data
(communication time is serial time)

```
#include <iostream>
#include <vector>
#include <thread>
#include <atomic>

static const int num_threads = 10;

void thread_code(const int tid, std::atomic<int>& result) {
    std::cout << tid << std::endl;

    result += 1;
}

int main(int argc, char** argv) {
    std::vector<std::thread> threads;
    std::atomic<int> result(0);

    threads.resize(num_threads);
    for (int i = 0; i < num_threads; ++i) {
        threads[i] = std::thread(thread_code, i, std::ref(result));
    }

    for (int i = 0; i < 10; ++i) {
        threads[i].join();
    }

    std::cout << "Result: " << result << std::endl;

    return EXIT_SUCCESS;
}
```

std::conditional_variable

allows sleep and wake-up of threads

```
#include <iostream>, #include <string>, #include <thread>, #include <mutex>, #include <condition_variable>

std::mutex m;
std::condition_variable cv;
std::string data;
bool ready = false;
bool processed = false;

void worker_thread()
{
    // Wait until main() sends data
    std::unique_lock<std::mutex> lk(m);
    cv.wait(lk, [] {return ready; });

    // after the wait, we own the lock.
    std::cout << "Worker thread is processing data\n";
    data += " after processing";

    // Send data back to main()
    processed = true;
    std::cout << "Worker thr. signals data processing completed\n";

    // Manual unlocking is done before notifying, to avoid
    // waking up the waiting thread only to block again (see
    // notify_one for details)
    lk.unlock();
    cv.notify_one();
}
```

```
int main()
{
    std::thread worker(worker_thread);

    data = "Example data";
    // send data to the worker thread
    {
        std::scoped_lock<std::mutex> lk(m);
        ready = true;
        std::cout << "main() signals data ready for processing\n";
    }
    cv.notify_one();

    // wait for the worker
    {
        std::unique_lock<std::mutex> lk(m);
        cv.wait(lk, [] {return processed; });
    }
    std::cout << "Back in main(), data = " << data << '\n';

    worker.join();
}
```

Thread-Safe Version of Container

- add lock to enforce exclusive access


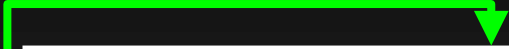
```
#include <deque>
#include <mutex>           // std::mutex, std::scoped_lock

template<typename T>
class synced_variable {
protected:
    std::mutex mux;
    T my_variable();

public:
    synced_variable() = default;
    synced_variable(const T&) = delete; //no copy-constructor

    // Guarded read
    const T& get() {
        std::scoped_lock lock(mux);
        return my_variable;
    }

    // Guarded write
    void set(const T& value) {
        std::scoped_lock lock(mux);
        my_variable = value;
    }
}
```



Thread-Safe Version of Container

- add lock to enforce exclusive access, and condition_variable to avoid busy-wait

```
#include <deque>
#include <mutex>           // std::mutex, std::scoped_lock
#include <condition_variable> // std::condition_variable

template<typename T>
class synced_deque {
protected:
    std::mutex mux;
    std::deque<T> de_queue;
    std::condition_variable cv_sleep;
    std::mutex mux_sleep;

public:
    synced_deque() = default;
    synced_deque(const synced_deque<T>&) = delete;
    virtual ~synced_deque() {
        clear();
    }

    // Returns and maintains item at front of Queue
    const T& front() {
        std::scoped_lock lock(mux);
        return de_queue.front();
    }

    // Removes and returns item from front of Queue
    T pop_front() {
        std::scoped_lock lock(mux);
        auto t = std::move(de_queue.front());
        de_queue.pop_front();
        return t;
    }
}
```

```
// Adds an item to back of Queue
void push_back(const T& item) {
    std::scoped_lock lock(mux);
    de_queue.emplace_back(std::move(item));

    std::unique_lock<std::mutex> ul(mux_sleep);
    cv_sleep.notify_one();
}

// Returns true if Queue has no items
bool empty() {
    std::scoped_lock lock(mux);
    return de_queue.empty();
}

void wait() {
    while (empty()) {
        std::unique_lock<std::mutex> ul(mux_sleep);
        cv_sleep.wait(ul);
    }
}

// const T& back();
// T pop_back()
// void push_front()
// const T& at();
// size_t size();
// ...
}
```

Thread Outside Class

VS.

Inside Class

```
#include <deque>
#include <iostream>
#include <thread>

void thread_code()
{

}

int main()
{
    std::thread worker(thread_code);

    worker.join();
}
```

```
#include <deque>
#include <iostream>
#include <thread>

class App {
public:
    App(); // default constructor
    void init(void);
    int run(void);
    void thread_code(void);
    ~App(); //default destructor
private:
};

void App::thread_code()
{

}

int App::run(void)
{
    // pass a pointer to current instance
    std::thread worker(&App::thread_code, this);
    ...
    worker.join();
}

int main()
{
    App app; // create instance

    return app.run();
}
```

Generative AI Example (fail)

Please create a simple C++ program with two threads: one thread is reading images using OpenCV and then send image to second thread. Second thread displays the image using OpenCV.

```
#include <iostream>
#include <thread>
#include <opencv2/opencv.hpp>

using namespace cv;

void readImage(Mat& image) {
    VideoCapture cap(0);
    while (true) {
        cap >> image;
        if (image.empty()) {
            std::cerr << "Error: no frame captured"
        }
        break;
    }
}

void displayImage(Mat& image) {
    while (true) {
        imshow("Display window", image);
        if (waitKey(1) >= 0) {
            break;
        }
    }
}

int main() {
    Mat image;
    std::thread t1(readImage, std::ref(image));
    std::thread t2(displayImage, std::ref(image));
    t1.join();
    t2.join();
    return 0;
}
```

Can you add synchronization between threads?

```
using namespace cv;

std::mutex mtx;

void readImage(Mat& image) {
    VideoCapture cap(0);
    while (true) {
        mtx.lock();
        cap >> image;
        if (image.empty()) {
            std::cerr << "Error: no frame captured"
        }
        break;
    }
    mtx.unlock();
}

void displayImage(Mat& image) {
    while (true) {
        mtx.lock();
        imshow("Display window", image);
        if (waitKey(1) >= 0) {
            break;
        }
        mtx.unlock();
    }
}

int main() {
    Mat image;
    std::thread t1(readImage, std::ref(image));
    std::thread t2(displayImage, std::ref(image));
    t1.join();
    t2.join();
    return 0;
}
```


Conclusion

- Concurrency is hard
- Use parallel algorithms when appropriate
- Avoid data races at all costs
 - Share less data
 - Lock guards (and mutexes)